

DATA621_Homework3

Jenny

November 4, 2018

```
install.packages("leaps") install.packages("lattice")
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
library(corrplot)
```

```
## corrplot 0.84 loaded
```

```
library(ggplot2)
```

```
library(reshape2)
```

```
library(leaps)
```

Introduction

In this homework, I will explore, analyze and model a data set containing information on crime for various neighborhoods of a major city. Each record has a response variable indicating whether or not the crime rate is above the median crime rate (1) or not (0).

My objective is to build a binary logistic regression model on the training data set to predict whether the neighborhood will be at risk for high crime levels. I will provide classifications and probabilities for the evaluation data set using my binary logistic regression model.

Read dataset

```
train_data <- read.csv("https://raw.githubusercontent.com/JennierJ/DATA621/master/Homework3/crime-train.csv")
eval_data <- read.csv("https://raw.githubusercontent.com/JennierJ/DATA621/master/Homework3/crime-evaluation.csv")
```

```
head(train_data)
```

```
##   zn indus chas   nox   rm   age   dis rad tax ptratio  black lstat medv
## 1  0 19.58    0 0.605 7.929 96.2 2.0459   5 403   14.7 369.30  3.70 50.0
## 2  0 19.58    1 0.871 5.403 100.0 1.3216   5 403   14.7 396.90 26.82 13.4
## 3  0 18.10    0 0.740 6.485 100.0 1.9784  24 666   20.2 386.73 18.85 15.4
## 4 30  4.93    0 0.428 6.393   7.8 7.0355   6 300   16.6 374.71  5.19 23.7
## 5  0  2.46    0 0.488 7.155 92.2 2.7006   3 193   17.8 394.12  4.82 37.9
## 6  0  8.56    0 0.520 6.781 71.3 2.8561   5 384   20.9 395.58  7.67 26.5
##   target
## 1      1
## 2      1
## 3      1
## 4      0
## 5      0
## 6      0
```

1. Data Exploration & Data Visualization

```
dim(train_data)
```

```
## [1] 466 14
```

```
summary(train_data)
```

```
##          zn          indus          chas          nox
## Min.   : 0.00   Min.   : 0.460   Min.   :0.00000   Min.   :0.3890
## 1st Qu.: 0.00   1st Qu.: 5.145   1st Qu.:0.00000   1st Qu.:0.4480
## Median : 0.00   Median : 9.690   Median :0.00000   Median :0.5380
## Mean   : 11.58   Mean   :11.105   Mean   :0.07082   Mean   :0.5543
## 3rd Qu.: 16.25   3rd Qu.:18.100   3rd Qu.:0.00000   3rd Qu.:0.6240
## Max.   :100.00   Max.   :27.740   Max.   :1.00000   Max.   :0.8710
##          rm          age          dis          rad
## Min.   :3.863   Min.   : 2.90   Min.   : 1.130   Min.   : 1.00
## 1st Qu.:5.887   1st Qu.: 43.88   1st Qu.: 2.101   1st Qu.: 4.00
## Median :6.210   Median : 77.15   Median : 3.191   Median : 5.00
## Mean   :6.291   Mean   : 68.37   Mean   : 3.796   Mean   : 9.53
## 3rd Qu.:6.630   3rd Qu.: 94.10   3rd Qu.: 5.215   3rd Qu.:24.00
## Max.   :8.780   Max.   :100.00   Max.   :12.127   Max.   :24.00
##          tax          ptratio          black          lstat
## Min.   :187.0   Min.   :12.6   Min.   : 0.32   Min.   : 1.730
## 1st Qu.:281.0   1st Qu.:16.9   1st Qu.:375.61   1st Qu.: 7.043
## Median :334.5   Median :18.9   Median :391.34   Median :11.350
## Mean   :409.5   Mean   :18.4   Mean   :357.12   Mean   :12.631
## 3rd Qu.:666.0   3rd Qu.:20.2   3rd Qu.:396.24   3rd Qu.:16.930
## Max.   :711.0   Max.   :22.0   Max.   :396.90   Max.   :37.970
##          medv          target
## Min.   : 5.00   Min.   :0.0000
## 1st Qu.:17.02   1st Qu.:0.0000
## Median :21.20   Median :0.0000
## Mean   :22.59   Mean   :0.4914
## 3rd Qu.:25.00   3rd Qu.:1.0000
## Max.   :50.00   Max.   :1.0000
```

The training dataset has 466 rows with 13 variables, with no NA values.

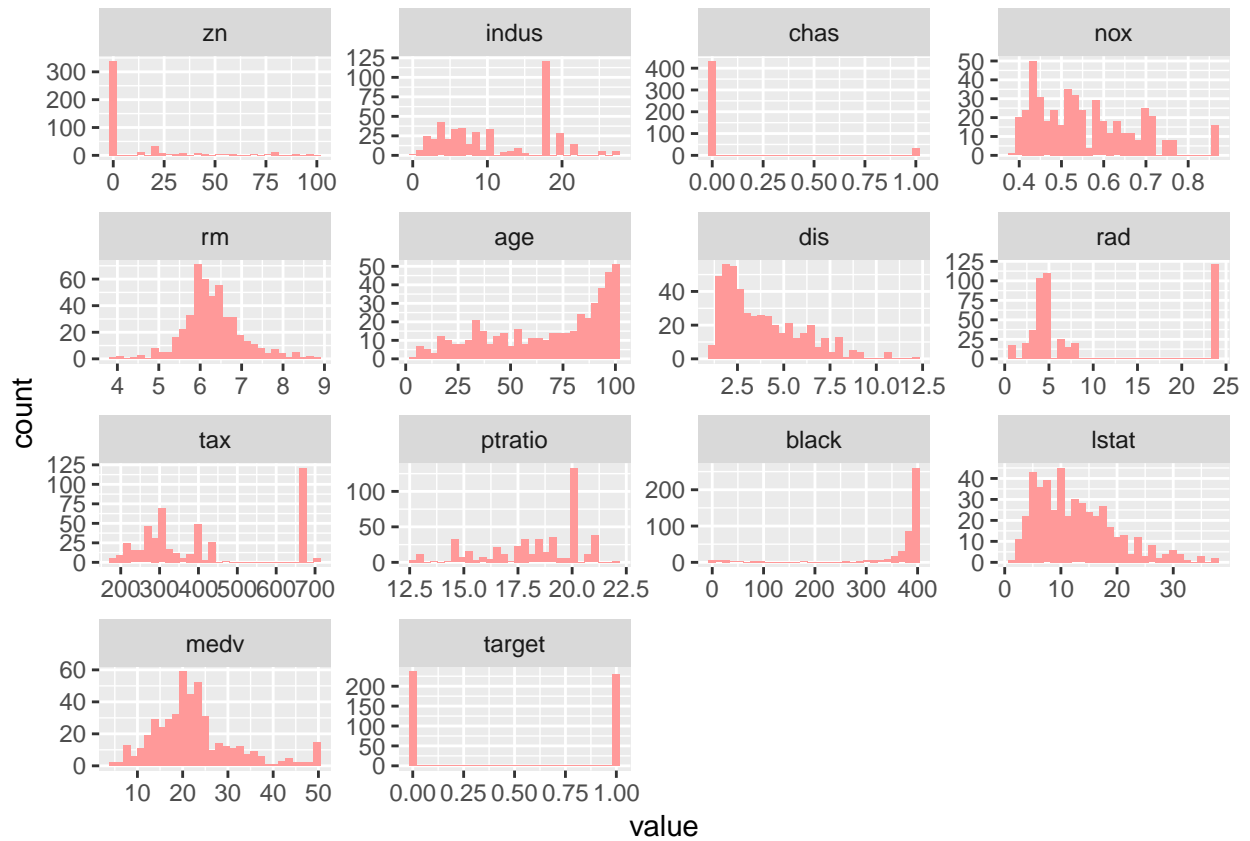
1.1. Histograms

```
v <- melt(train_data)
```

```
## No id variables; using all as measure variables
```

```
ggplot(v, aes(value)) + geom_histogram(fill = "#FF9999") + facet_wrap(~variable,
                                                                    scales = "free")
```

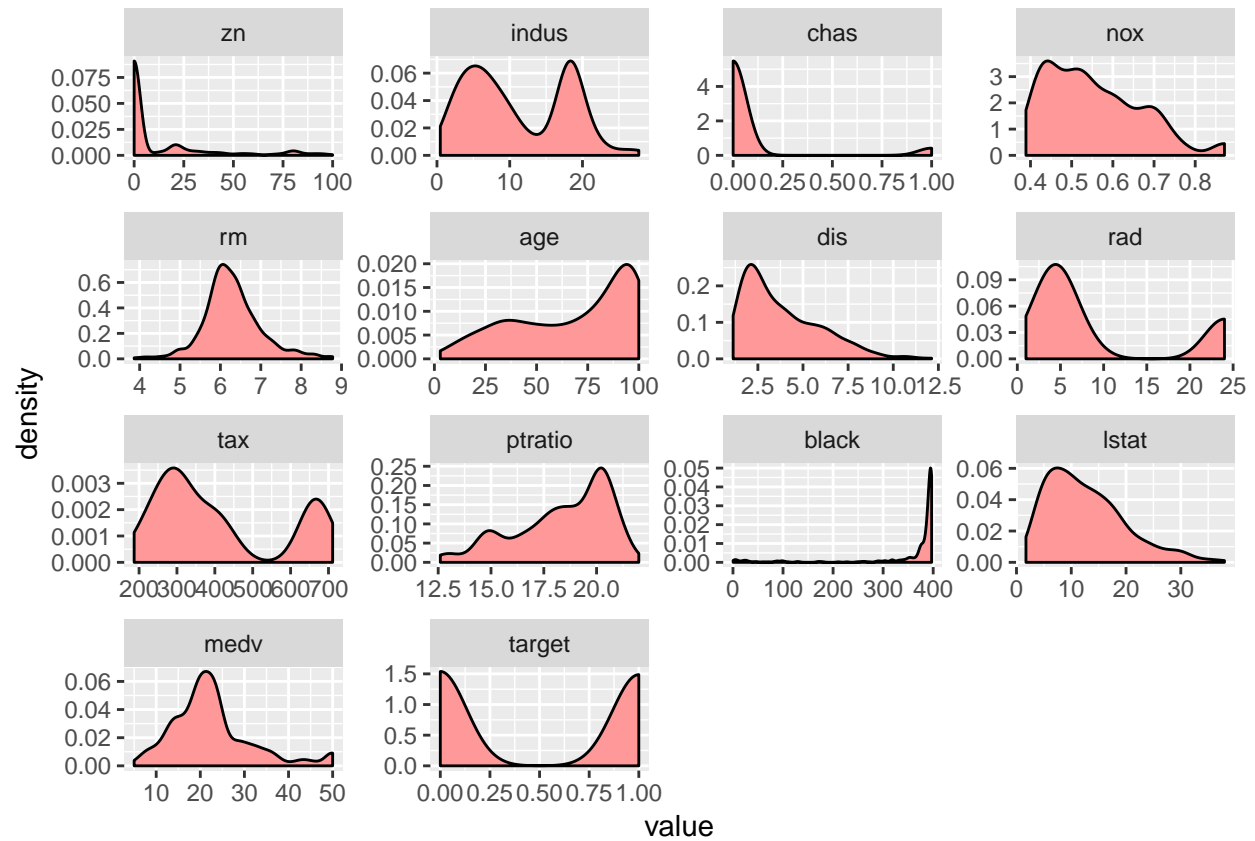
```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



From the histograms above, it looks like that the variable zn, dis and age are skewed.


1.2 Density

```
ggplot(v, aes(value)) + geom_density(fill = "#FF9999") + facet_wrap(~variable,
                                                                    scales = "free")
```



1.3 Correlation

```
corrplot(cor(train_data), method = "color")
```



The correlation plot has shown that how variable in the dataset are related to each other.

2. Data Preparation

2.1. Data Transformation

I would like to do log transformation to medv and rad.

```
train_data$lgmedv <- log(train_data$medv)
train_data$lggrad <- log(train_data$rad)
```

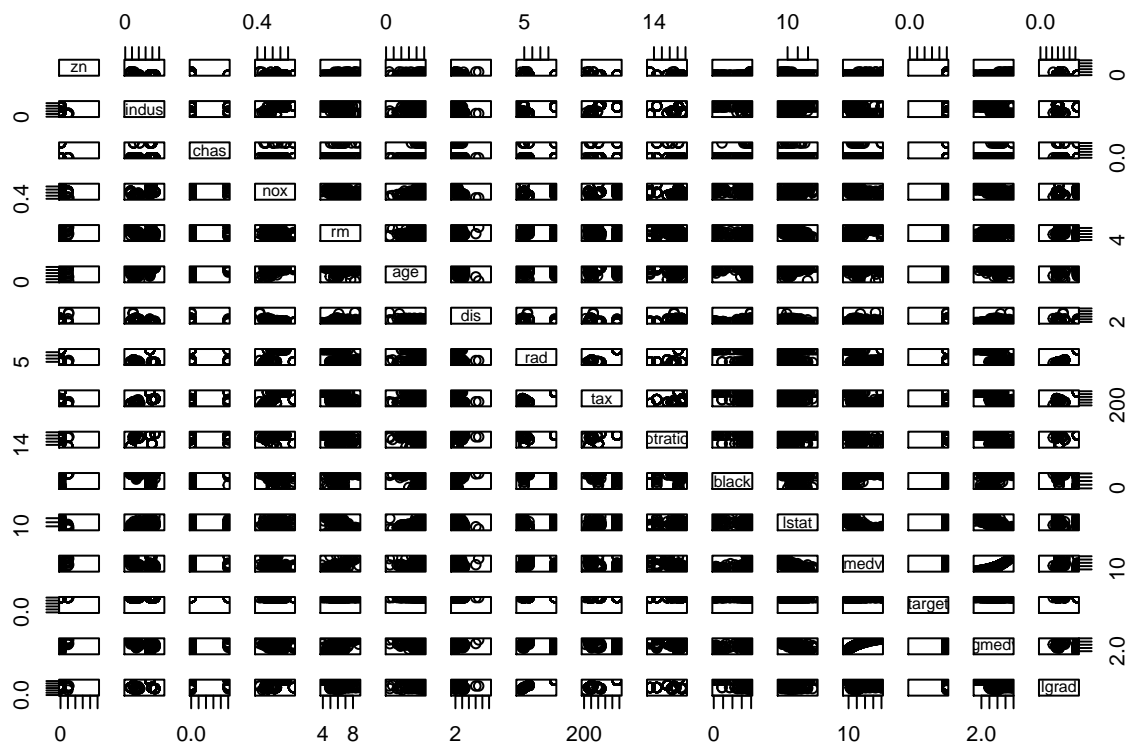
```
head(train_data)
```

```
##   zn indus chas   nox    rm   age    dis rad tax ptratio  black  lstat medv
## 1  0 19.58   0 0.605 7.929 96.2 2.0459  5 403    14.7 369.30  3.70 50.0
## 2  0 19.58   1 0.871 5.403 100.0 1.3216  5 403    14.7 396.90 26.82 13.4
## 3  0 18.10   0 0.740 6.485 100.0 1.9784 24 666    20.2 386.73 18.85 15.4
## 4 30  4.93   0 0.428 6.393  7.8 7.0355  6 300    16.6 374.71  5.19 23.7
## 5  0  2.46   0 0.488 7.155 92.2 2.7006  3 193    17.8 394.12  4.82 37.9
## 6  0  8.56   0 0.520 6.781 71.3 2.8561  5 384    20.9 395.58  7.67 26.5
##   target  lgmedv  lggrad
## 1      1  3.912023 1.609438
## 2      1  2.595255 1.609438
## 3      1  2.734368 3.178054
## 4      0  3.165475 1.791759
```

```
## 5      0 3.634951 1.098612
## 6      0 3.277145 1.609438
```

3. Build Model

```
pairs(train_data, col = train_data$target)
```



```
#### I would like to use two methods to create models.
```

3.1 Logistic Regression Backward Selection

```
glm.fit = glm(train_data$target ~ train_data$zn + train_data$indus + train_data$chas + train_data$nox
              + train_data$rm + train_data$age + train_data$dis + train_data$lgrad + train_data$tax +
              train_data$ptratio + train_data$black + train_data$lstat + train_data$lmedv, data = tr
              family = binomial)
summary(glm.fit)
```

```
##
## Call:
## glm(formula = train_data$target ~ train_data$zn + train_data$indus +
##      train_data$chas + train_data$nox + train_data$rm + train_data$age +
##      train_data$dis + train_data$lgrad + train_data$tax + train_data$ptratio +
##      train_data$black + train_data$lstat + train_data$lmedv,
##      family = binomial, data = train_data)
##
```

```
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.5718  -0.1243  -0.0008   0.0671   3.4238
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -48.528140    9.439591  -5.141 2.73e-07 ***
## train_data$zn     -0.055872    0.032045  -1.744  0.08124 .
## train_data$indus  -0.039320    0.050844  -0.773  0.43932
## train_data$chas    0.830527    0.764727   1.086  0.27746
## train_data$nox    46.955469    7.912562   5.934 2.95e-09 ***
## train_data$rm      0.131597    0.623713   0.211  0.83290
## train_data$age     0.028696    0.013133   2.185  0.02888 *
## train_data$dis     0.648672    0.225915   2.871  0.00409 **
## train_data$lgrad   3.391947    0.771016   4.399 1.09e-05 ***
## train_data$tax    -0.007772    0.003508  -2.215  0.02674 *
## train_data$prratio 0.441686    0.134305   3.289  0.00101 **
## train_data$black  -0.013830    0.007319  -1.890  0.05880 .
## train_data$lstat   0.071499    0.055665   1.284  0.19898
## train_data$lgmedv  3.823376    1.760816   2.171  0.02990 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 189.17  on 452  degrees of freedom
## AIC: 217.17
##
## Number of Fisher Scoring iterations: 8
```

In the summary of model, I would like to remove the variable with a P value higher than 0.05,
- variable rm

```
glm.fit1 = glm(train_data$target ~ train_data$zn + train_data$indus + train_data$chas + train_data$nox
              + train_data$age + train_data$dis + train_data$lgrad + train_data$tax +
              train_data$prratio + train_data$black + train_data$lstat + train_data$lgmedv, data = tr
              family = binomial)
summary(glm.fit1)
```

```
##
## Call:
## glm(formula = train_data$target ~ train_data$zn + train_data$indus +
##      train_data$chas + train_data$nox + train_data$age + train_data$dis +
##      train_data$lgrad + train_data$tax + train_data$prratio +
##      train_data$black + train_data$lstat + train_data$lgmedv,
##      family = binomial, data = train_data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.5751  -0.1227  -0.0008   0.0678   3.4264
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
```

```
## (Intercept)          -48.916647    9.311444   -5.253 1.49e-07 ***
## train_data$zn         -0.054929    0.031674   -1.734 0.082879 .
## train_data$indus      -0.040254    0.050722   -0.794 0.427415
## train_data$chas        0.822779    0.763423    1.078 0.281145
## train_data$nox        47.222154    7.839626    6.024 1.71e-09 ***
## train_data$age         0.030069    0.011446    2.627 0.008614 **
## train_data$dis         0.658434    0.221763    2.969 0.002987 **
## train_data$lgrad       3.407807    0.769793    4.427 9.56e-06 ***
## train_data$tax        -0.007741    0.003517   -2.201 0.027725 *
## train_data$ptratio     0.451348    0.126940    3.556 0.000377 ***
## train_data$black      -0.013966    0.007298   -1.914 0.055656 .
## train_data$lstat       0.068638    0.054105    1.269 0.204578
## train_data$lgmedv      4.085496    1.253737    3.259 0.001119 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 189.21  on 453  degrees of freedom
## AIC: 215.21
##
## Number of Fisher Scoring iterations: 8
```

- variable indus

```
glm.fit2 = glm(train_data$target ~ train_data$zn + train_data$chas + train_data$nox
               + train_data$age + train_data$dis + train_data$lgrad + train_data$tax +
               train_data$ptratio + train_data$black + train_data$lstat + train_data$lgmedv, data = t,
               family = binomial)
summary(glm.fit2)
```

```
##
## Call:
## glm(formula = train_data$target ~ train_data$zn + train_data$chas +
##      train_data$nox + train_data$age + train_data$dis + train_data$lgrad +
##      train_data$tax + train_data$ptratio + train_data$black +
##      train_data$lstat + train_data$lgmedv, family = binomial,
##      data = train_data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.4950  -0.1260  -0.0007   0.0642   3.3997
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -47.459673    9.074063  -5.230 1.69e-07 ***
## train_data$zn   -0.054303    0.030963  -1.754 0.07947 .
## train_data$chas    0.676716    0.742654   0.911 0.36218
## train_data$nox   44.662991    6.973847   6.404 1.51e-10 ***
## train_data$age     0.029758    0.011328   2.627 0.00862 **
## train_data$dis     0.639033    0.219027   2.918 0.00353 **
## train_data$lgrad    3.660451    0.726530   5.038 4.70e-07 ***
## train_data$tax    -0.009000    0.003116  -2.888 0.00388 **
```



```
## train_data$ptratio    0.440650    0.126604    3.481  0.00050 ***
## train_data$black      -0.013482    0.007108   -1.897  0.05785 .
## train_data$lstat      0.065282    0.053675    1.216  0.22389
## train_data$lgmedv     3.982421    1.246465    3.195  0.00140 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 189.86  on 454  degrees of freedom
## AIC: 213.86
##
## Number of Fisher Scoring iterations: 8
```

- variable chas

```
glm.fit3 = glm(train_data$target ~ train_data$zn + train_data$nox
               + train_data$age + train_data$dis + train_data$lgrad + train_data$tax +
               train_data$ptratio + train_data$black + train_data$lstat + train_data$lgmedv, data = t,
               family = binomial)
summary(glm.fit3)
```

```
##
## Call:
## glm(formula = train_data$target ~ train_data$zn + train_data$nox +
##      train_data$age + train_data$dis + train_data$lgrad + train_data$tax +
##      train_data$ptratio + train_data$black + train_data$lstat +
##      train_data$lgmedv, family = binomial, data = train_data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.4745  -0.1274  -0.0005   0.0646   3.4076
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -47.124386    9.046356  -5.209 1.90e-07 ***
## train_data$zn    -0.059665    0.031116  -1.917 0.055177 .
## train_data$nox    43.877043    6.882902   6.375 1.83e-10 ***
## train_data$age     0.030488    0.011270   2.705 0.006825 **
## train_data$dis     0.626269    0.217594   2.878 0.004000 **
## train_data$lgrad   3.787325    0.729067   5.195 2.05e-07 ***
## train_data$tax    -0.009436    0.003149  -2.997 0.002729 **
## train_data$ptratio  0.421373    0.123778   3.404 0.000663 ***
## train_data$black  -0.013292    0.007070  -1.880 0.060099 .
## train_data$lstat   0.074672    0.052362   1.426 0.153847
## train_data$lgmedv  4.066484    1.243289   3.271 0.001073 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 190.71  on 455  degrees of freedom
```

```
## AIC: 212.71
##
## Number of Fisher Scoring iterations: 8
```

- variable zn

```
glm.fit4 = glm(train_data$target ~ train_data$nox
               + train_data$age + train_data$dis + train_data$lgrad + train_data$tax +
               train_data$ptratio + train_data$black + train_data$lstat + train_data$lgmedv, data = t,
               family = binomial)
summary(glm.fit4)
```

```
##
## Call:
## glm(formula = train_data$target ~ train_data$nox + train_data$age +
##      train_data$dis + train_data$lgrad + train_data$tax + train_data$ptratio +
##      train_data$black + train_data$lstat + train_data$lgmedv,
##      family = binomial, data = train_data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.60483  -0.18125  -0.00248   0.05699   3.12228
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -45.638033    8.857960  -5.152 2.57e-07 ***
## train_data$nox    43.755818    6.911380   6.331 2.44e-10 ***
## train_data$age     0.029892    0.011115   2.689 0.007158 **
## train_data$dis     0.428557    0.178563   2.400 0.016394 *
## train_data$lgrad    3.735967    0.687829   5.432 5.59e-08 ***
## train_data$tax    -0.009725    0.002947  -3.300 0.000967 ***
## train_data$ptratio  0.475082    0.121170   3.921 8.83e-05 ***
## train_data$black  -0.013822    0.007300  -1.893 0.058302 .
## train_data$lstat    0.065994    0.053095   1.243 0.213884
## train_data$lgmedv   3.625056    1.216385   2.980 0.002881 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 195.50  on 456  degrees of freedom
## AIC: 215.5
##
## Number of Fisher Scoring iterations: 8
```

- variable black

```
glm.fit5 = glm(train_data$target ~ train_data$nox
               + train_data$age + train_data$dis + train_data$lgrad + train_data$tax +
               train_data$ptratio + train_data$lstat + train_data$lgmedv, data = train_data,
               family = binomial)
summary(glm.fit5)
```

```
##
## Call:
## glm(formula = train_data$target ~ train_data$nox + train_data$age +
##      train_data$dis + train_data$lgrad + train_data$tax + train_data$ptratio +
##      train_data$lstat + train_data$lgmedv, family = binomial,
##      data = train_data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.07722  -0.18052  -0.00237   0.07599   3.11297
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -48.422975     8.481067  -5.710 1.13e-08 ***
## train_data$nox     43.923656     6.895490   6.370 1.89e-10 ***
## train_data$age      0.029837     0.011201   2.664 0.007727 **
## train_data$dis      0.394632     0.177807   2.219 0.026457 *
## train_data$lgrad    3.782159     0.664560   5.691 1.26e-08 ***
## train_data$tax     -0.009417     0.002844  -3.311 0.000930 ***
## train_data$ptratio  0.440743     0.116615   3.779 0.000157 ***
## train_data$lstat    0.053130     0.053254   0.998 0.318441
## train_data$lgmedv   3.035936     1.161650   2.613 0.008963 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 201.51  on 457  degrees of freedom
## AIC: 219.51
##
## Number of Fisher Scoring iterations: 8
```

- variable lstat

```
glm.fit6 = glm(target ~ nox
               + age + dis + lgrad + tax +
               ptratio + lgmedv, data = train_data,
               family = binomial)
summary(glm.fit6)
```

```
##
## Call:
## glm(formula = target ~ nox + age + dis + lgrad + tax + ptratio +
##      lgmedv, family = binomial, data = train_data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.02369  -0.18216  -0.00268   0.09140   3.13756
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -45.173968     7.682267  -5.880 4.1e-09 ***
## nox          43.568175     6.822642   6.386 1.7e-10 ***
```

```
## age          0.032602    0.010887    2.995 0.002748 **
## dis          0.383970    0.177580    2.162 0.030600 *
## lgrad        3.739570    0.654334    5.715 1.1e-08 ***
## tax         -0.009386    0.002837   -3.309 0.000936 ***
## ptratio      0.421167    0.114560    3.676 0.000237 ***
## lgmedv       2.335450    0.918850    2.542 0.011031 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 202.51  on 458  degrees of freedom
## AIC: 218.51
##
## Number of Fisher Scoring iterations: 8
```

Now that the model glm.fit 6 are having variables left with P value lower than 0.05.

3.2 Lead Model Selection

The leaps packages is helping to generate all subset regression models.

```
regfit.full = regsubsets(target~., data = train_data, nvmax = 15)
reg.summary = summary(regfit.full)
reg.summary
```

```
## Subset selection object
## Call: regsubsets.formula(target ~ ., data = train_data, nvmax = 15)
## 15 Variables (and intercept)
##           Forced in Forced out
## zn          FALSE      FALSE
## indus       FALSE      FALSE
## chas        FALSE      FALSE
## nox         FALSE      FALSE
## rm          FALSE      FALSE
## age         FALSE      FALSE
## dis         FALSE      FALSE
## rad         FALSE      FALSE
## tax         FALSE      FALSE
## ptratio     FALSE      FALSE
## black       FALSE      FALSE
## lstat       FALSE      FALSE
## medv        FALSE      FALSE
## lgmedv      FALSE      FALSE
## lgrad       FALSE      FALSE
## 1 subsets of each size up to 15
## Selection Algorithm: exhaustive
##           zn indus chas nox rm age dis rad tax ptratio black lstat medv
## 1 ( 1 ) " " " " " " "*" " " " " " " " " " " " " " " " "
## 2 ( 1 ) " " " " " " "*" " " " " " " " " " " " " " " "
## 3 ( 1 ) " " " " " " "*" " " "*" " " " " " " " " " " "
## 4 ( 1 ) " " " " " " "*" " " "*" " " " " " " " " " " "*"
```

```

## 5 ( 1 ) " " " " " " "*" " " "*" " " " " " " "*" " " " " "*"
## 6 ( 1 ) " " " " " " "*" " " "*" " " " " " " "*" "*" " " "*"
## 7 ( 1 ) " " " " " " "*" " " "*" " " " " " " "*" "*" "*" "*"
## 8 ( 1 ) " " " " " " "*" "*" "*" " " " " " " "*" "*" "*" "*"
## 9 ( 1 ) " " " " " " "*" "*" "*" " " " " " " "*" "*" "*" "*"
## 10 ( 1 ) "*" " " " " "*" "*" "*" "*" " " " " " " "*" "*" "*" "*"
## 11 ( 1 ) " " "*" " " " "*" "*" "*" " " "*" "*" "*" "*" "*" "*" "*"
## 12 ( 1 ) " " "*" " " " "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*"
## 13 ( 1 ) "*" "*" " " " "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*"
## 14 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*"
## 15 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*"
##          lgmedv lgrad
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " "*"
## 3 ( 1 ) " " "*"
## 4 ( 1 ) " " "*"
## 5 ( 1 ) " " "*"
## 6 ( 1 ) " " "*"
## 7 ( 1 ) " " "*"
## 8 ( 1 ) " " "*"
## 9 ( 1 ) " " "*"
## 10 ( 1 ) " " "*"
## 11 ( 1 ) " " "*"
## 12 ( 1 ) " " "*"
## 13 ( 1 ) " " "*"
## 14 ( 1 ) " " "*"
## 15 ( 1 ) "*" "*"

names(reg.summary)


## [1] "which" "rsq" "rss" "adjr2" "cp" "bic" "outmat" "obj"

plot(reg.summary$cp, xlab = "Number of Variables", ylab = "Cp")
which.min(reg.summary$cp)

## [1] 5

points(5, reg.summary$cp[5], pch = 20, col = "red")

```



```
glm.fit7 = glm(train_data$target ~ train_data$nox
               + train_data$age + train_data$ptratio + train_data$medv + train_data$lgrad, data = train_data,
               family = binomial)
summary(glm.fit7)
```

```
##
## Call:
## glm(formula = train_data$target ~ train_data$nox + train_data$age +
##      train_data$ptratio + train_data$medv + train_data$lgrad,
##      family = binomial, data = train_data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.03040  -0.27870  -0.01536   0.08021   2.86196
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -27.254311   3.773829  -7.222 5.13e-13 ***
## train_data$nox    25.669907   4.098389   6.263 3.77e-10 ***
## train_data$age     0.019977   0.009248   2.160 0.030751 *
## train_data$ptratio  0.307056   0.100031   3.070 0.002143 **
## train_data$medv    0.094301   0.028525   3.306 0.000947 ***
## train_data$lgrad   2.478406   0.495238   5.004 5.60e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 645.88 on 465 degrees of freedom
## Residual deviance: 222.88 on 460 degrees of freedom
## AIC: 234.88
##
## Number of Fisher Scoring iterations: 7
```

The model `glm.fit7` is the “best” model that with lowest CP and picked by leap package

4. Model Selection

Let's see the factor of model `glm.fit6`

```
summary(glm.fit6)
```

```
##
## Call:
## glm(formula = target ~ nox + age + dis + lgrad + tax + ptratio +
##      lgmedv, family = binomial, data = train_data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.02369  -0.18216  -0.00268   0.09140   3.13756
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -45.173968    7.682267  -5.880  4.1e-09 ***
## nox          43.568175    6.822642   6.386  1.7e-10 ***
## age           0.032602    0.010887   2.995  0.002748 **
## dis           0.383970    0.177580   2.162  0.030600 *
## lgrad         3.739570    0.654334   5.715  1.1e-08 ***
## tax          -0.009386    0.002837  -3.309  0.000936 ***
## ptratio       0.421167    0.114560   3.676  0.000237 ***
## lgmedv        2.335450    0.918850   2.542  0.011031 *
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## (Dispersion parameter for binomial family taken to be 1)
##
```

```
## Null deviance: 645.88 on 465 degrees of freedom
## Residual deviance: 202.51 on 458 degrees of freedom
## AIC: 218.51
##
```

```
## Number of Fisher Scoring iterations: 8
```

```
glm.probs = predict(glm.fit6, data = train_data, type = "response")
```

```
Matrix6 <- confusionMatrix(data = factor(ifelse(glm.probs > 0.5, 1, 0)), reference = factor(train_data$
      positive = "1"))
```

```
Matrix6
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction  0   1
##           0 219  24
##           1  18 205
##
##           Accuracy : 0.9099
##           95% CI : (0.8801, 0.9343)
##           No Information Rate : 0.5086
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.8196
## Mcnemar's Test P-Value : 0.4404
##
##           Sensitivity : 0.8952
##           Specificity : 0.9241
##           Pos Pred Value : 0.9193
##           Neg Pred Value : 0.9012
##           Prevalence : 0.4914
##           Detection Rate : 0.4399
##           Detection Prevalence : 0.4785
##           Balanced Accuracy : 0.9096
##
##           'Positive' Class : 1
##
```

And the factor of model glm.fit7

```
summary(glm.fit7)
```

```
##
## Call:
## glm(formula = train_data$target ~ train_data$nox + train_data$age +
##      train_data$ptratio + train_data$medv + train_data$lgrad,
##      family = binomial, data = train_data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.03040  -0.27870  -0.01536   0.08021   2.86196
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -27.254311   3.773829  -7.222 5.13e-13 ***
## train_data$nox    25.669907   4.098389   6.263 3.77e-10 ***
## train_data$age     0.019977   0.009248   2.160 0.030751 *
## train_data$ptratio  0.307056   0.100031   3.070 0.002143 **
## train_data$medv    0.094301   0.028525   3.306 0.000947 ***
## train_data$lgrad   2.478406   0.495238   5.004 5.60e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
```



```
## Null deviance: 645.88 on 465 degrees of freedom
## Residual deviance: 222.88 on 460 degrees of freedom
## AIC: 234.88
##
## Number of Fisher Scoring iterations: 7
glm.probs = predict(glm.fit7, data = train_data, type = "response")

Matrix7 <- confusionMatrix(data = factor(ifelse(glm.probs > 0.5, 1, 0)), reference = factor(train_data$
positive = "1"))
Matrix7

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0 209  36
##           1   28 193
##
##           Accuracy : 0.8627
##           95% CI : (0.828, 0.8926)
##       No Information Rate : 0.5086
##       P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.7251
##  Mcnemar's Test P-Value : 0.3816
##
##           Sensitivity : 0.8428
##           Specificity : 0.8819
##       Pos Pred Value : 0.8733
##       Neg Pred Value : 0.8531
##           Prevalence : 0.4914
##       Detection Rate : 0.4142
##       Detection Prevalence : 0.4742
##       Balanced Accuracy : 0.8623
##
##       'Positive' Class : 1
##
```

Using model glm.fit6 to predict the target with evaluation_data

```
eval_data$lgrad <- log(eval_data$rad)
eval_data$lgmedv <- log(eval_data$medv)
glm.probs.pred = predict(glm.fit6, newdata = eval_data, type = "response")
predtarget <- ifelse(glm.probs.pred > 0.5, 1, 0)
eval_data$pred <- predtarget
eval_data
```

```
##      zn indus chas   nox    rm   age    dis rad tax ptratio  black lstat
## 1    0  7.07    0 0.469 7.185 61.1 4.9671  2 242    17.8 392.83  4.03
## 2    0  8.14    0 0.538 6.096 84.5 4.4619  4 307    21.0 380.02 10.26
## 3    0  8.14    0 0.538 6.495 94.4 4.4547  4 307    21.0 387.94 12.80
## 4    0  8.14    0 0.538 5.950 82.0 3.9900  4 307    21.0 232.60 27.71
## 5    0  5.96    0 0.499 5.850 41.5 3.9342  5 279    19.2 396.90  8.77
```

## 6	25	5.13	0	0.453	5.741	66.2	7.2254	8	284	19.7	395.11	13.15
## 7	25	5.13	0	0.453	5.966	93.4	6.8185	8	284	19.7	378.08	14.44
## 8	0	4.49	0	0.449	6.630	56.1	4.4377	3	247	18.5	392.30	6.53
## 9	0	4.49	0	0.449	6.121	56.8	3.7476	3	247	18.5	395.15	8.44
## 10	0	2.89	0	0.445	6.163	69.6	3.4952	2	276	18.0	391.83	11.34
## 11	0	25.65	0	0.581	5.856	97.0	1.9444	2	188	19.1	370.31	25.41
## 12	0	25.65	0	0.581	5.613	95.6	1.7572	2	188	19.1	359.29	27.26
## 13	0	21.89	0	0.624	5.637	94.7	1.9799	4	437	21.2	396.90	18.34
## 14	0	19.58	0	0.605	6.101	93.0	2.2834	5	403	14.7	240.16	9.81
## 15	0	19.58	0	0.605	5.880	97.3	2.3887	5	403	14.7	348.13	12.03
## 16	0	10.59	1	0.489	5.960	92.1	3.8771	4	277	18.6	393.25	17.27
## 17	0	6.20	0	0.504	6.552	21.4	3.3751	8	307	17.4	380.34	3.76
## 18	0	6.20	0	0.507	8.247	70.4	3.6519	8	307	17.4	378.95	3.95
## 19	22	5.86	0	0.431	6.957	6.8	8.9067	7	330	19.1	386.09	3.53
## 20	90	2.97	0	0.400	7.088	20.8	7.3073	1	285	15.3	394.72	7.85
## 21	80	1.76	0	0.385	6.230	31.5	9.0892	1	241	18.2	341.60	12.93
## 22	33	2.18	0	0.472	6.616	58.1	3.3700	7	222	18.4	393.36	8.93
## 23	0	9.90	0	0.544	6.122	52.8	2.6403	4	304	18.4	396.90	5.98
## 24	0	7.38	0	0.493	6.415	40.1	4.7211	5	287	19.6	396.90	6.12
## 25	0	7.38	0	0.493	6.312	28.9	5.4159	5	287	19.6	396.90	6.15
## 26	0	5.19	0	0.515	5.895	59.6	5.6150	5	224	20.2	394.81	10.56
## 27	80	2.01	0	0.435	6.635	29.7	8.3440	4	280	17.0	390.94	5.99
## 28	0	18.10	0	0.718	3.561	87.9	1.6132	24	666	20.2	354.70	7.12
## 29	0	18.10	1	0.631	7.016	97.5	1.2024	24	666	20.2	392.05	2.96
## 30	0	18.10	0	0.584	6.348	86.1	2.0527	24	666	20.2	83.45	17.64
## 31	0	18.10	0	0.740	5.935	87.9	1.8206	24	666	20.2	68.95	34.02
## 32	0	18.10	0	0.740	5.627	93.9	1.8172	24	666	20.2	396.90	22.88
## 33	0	18.10	0	0.740	5.818	92.4	1.8662	24	666	20.2	391.45	22.11
## 34	0	18.10	0	0.740	6.219	100.0	2.0048	24	666	20.2	395.69	16.59
## 35	0	18.10	0	0.740	5.854	96.6	1.8956	24	666	20.2	240.52	23.79
## 36	0	18.10	0	0.713	6.525	86.5	2.4358	24	666	20.2	50.92	18.13
## 37	0	18.10	0	0.713	6.376	88.4	2.5671	24	666	20.2	391.43	14.65
## 38	0	18.10	0	0.655	6.209	65.4	2.9634	24	666	20.2	396.90	13.22
## 39	0	9.69	0	0.585	5.794	70.6	2.8927	6	391	19.2	396.90	14.10
## 40	0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0	396.90	5.64
##	medv		lgrad	lgmedv	pred							
## 1	34.7	0.6931472	3.546740		0							
## 2	18.2	1.3862944	2.901422		1							
## 3	18.4	1.3862944	2.912351		1							
## 4	13.2	1.3862944	2.580217		0							
## 5	21.0	1.6094379	3.044522		0							
## 6	18.7	2.0794415	2.928524		0							
## 7	16.0	2.0794415	2.772589		1							
## 8	26.6	1.0986123	3.280911		0							
## 9	22.2	1.0986123	3.100092		0							
## 10	21.4	0.6931472	3.063391		0							
## 11	17.3	0.6931472	2.850707		0							
## 12	15.7	0.6931472	2.753661		0							
## 13	14.3	1.3862944	2.660260		1							
## 14	25.0	1.6094379	3.218876		1							
## 15	19.1	1.6094379	2.949688		1							
## 16	21.7	1.3862944	3.077312		0							
## 17	31.5	2.0794415	3.449988		0							
## 18	48.3	2.0794415	3.877432		1							

```
## 19 29.6 1.9459101 3.387774 0
## 20 32.2 0.0000000 3.471966 0
## 21 20.1 0.0000000 3.000720 0
## 22 28.4 1.9459101 3.346389 0
## 23 22.1 1.3862944 3.095578 0
## 24 25.0 1.6094379 3.218876 0
## 25 23.0 1.6094379 3.135494 0
## 26 18.5 1.6094379 2.917771 1
## 27 24.5 1.3862944 3.198673 0
## 28 27.5 3.1780538 3.314186 1
## 29 50.0 3.1780538 3.912023 1
## 30 14.5 3.1780538 2.674149 1
## 31 8.4 3.1780538 2.128232 1
## 32 12.8 3.1780538 2.549445 1
## 33 10.5 3.1780538 2.351375 1
## 34 18.4 3.1780538 2.912351 1
## 35 10.8 3.1780538 2.379546 1
## 36 14.1 3.1780538 2.646175 1
## 37 17.7 3.1780538 2.873565 1
## 38 21.4 3.1780538 3.063391 1
## 39 18.3 1.7917595 2.906901 1
## 40 23.9 0.0000000 3.173878 0
```

```
table(predtarget)
```

```
## predtarget
## 0 1
## 20 20
```

```
summary(glm.probs.pred)
```

```
##      Min.    1st Qu.    Median      Mean    3rd Qu.      Max.
## 0.0000042 0.1394000 0.4815000 0.5164000 0.9894000 1.0000000
```